

SASI SDK Administrator Guide

Public Website Edition

Version: 1.3.1 Public

Last Updated: February 2026

Audience: Administrators, Compliance Officers, Clinical and Safety Staff

Note: This is the public edition. It is intentionally detailed about what SASI does and how teams operate it, while withholding the exact implementation specifics that would enable replication. Customers receive the full administrator manual plus the integration and audit specification pack after purchase.

Table of Contents

1. Introduction
2. Platform Overview
3. Dashboard and Navigation
4. Configuration and Controls
5. Reports and Analytics
6. Governance, Audits, and Change Management
7. Performance Monitoring
8. Operational Modes
9. Quick Reference
10. Support and Resources

1. Introduction

What is SASI

SASI, Symbolic AI Safety Intelligence, is safety middleware that analyzes user messages before they reach your AI application. It sits between your users and your AI system to detect risk, protect sensitive information, and generate compliance grade records.

What SASI does not do

1. SASI does not generate AI responses
2. SASI does not replace human oversight
3. SASI does not make clinical or medical decisions
4. SASI does not require you to rewrite your product logic to function

What SASI does

1. Analyzes messages for safety risk patterns, including crisis signals and emotional distress
2. Detects and redacts personally identifiable information, with stricter handling in regulated contexts
3. Recommends actions your product can take, such as continue, escalate, or present crisis resources
4. Creates audit trails designed for regulated environments and enterprise governance
5. Monitors system performance and safety behavior over time, including drift detection

Key concepts

Operational Mode

SASI runs in a mode aligned to your use case. Each mode changes sensitivity, privacy posture, reporting depth, and compliance expectations.

Risk Levels

SASI assigns a risk level per message. Levels are ordered from normal to urgent, and are designed to be readable by product teams, safety teams, and auditors.

Recommended Actions

SASI provides a recommended action per message. Your application chooses how to implement actions, but SASI ensures recommendations are consistent, explainable, and auditable.

2. Platform Overview

SASI is typically deployed as an SDK plus an administrative console.

1. The SDK performs message analysis, redaction, and safety decisions
2. The console exposes configuration, reporting, governance controls, and monitoring views
3. Events can be pushed into your tooling via notifications and integrations

Public edition note

Exact data schemas, integration field maps, payload shapes, and enforcement rules are provided in the customer implementation pack.

3. Dashboard and Navigation

Common dashboard areas

1. **Overview**
High level safety and privacy metrics, recent activity, and system health.
2. **Settings**
Operational mode, privacy controls, safety sensitivity, notification routing, and access controls.
3. **Reports**
Audit records, incident reports, privacy reports, exports, and analytics views.
4. **Monitoring**
Real time performance, active escalations, alerting, and drift monitoring.
5. **Incidents**
Crisis detections, escalations, and investigation workflows.
6. **Compliance**
Regulatory oriented exports, retention settings, and audit readiness tools.

Metrics you will typically see

Crisis detection rate

A small percentage of messages that represent potential crisis content. The expected range varies by mode.

PII detection rate

The percentage of messages containing sensitive identifiers. Higher is normal in healthcare and youth contexts.

Processing time

The time SASI spends per message. This depends on tier and mode and is usually designed to be low latency for interactive chat.

False positive indicators

Signals that safe messages are being over flagged. These are monitored to protect user experience while maintaining safety.

Adversarial attempt rate

Signals of manipulation attempts against your AI system. This should usually remain very low, but it varies by domain.

4. Configuration and Controls

4.1 Operational mode

Operational mode is the most important setting. It defines the baseline safety posture and compliance behavior.

Changing mode affects

1. Crisis sensitivity
2. PII detection scope and redaction strictness
3. Audit depth and retention expectations
4. Available safety and governance features

Safety locks

In regulated and high risk modes, certain protections cannot be disabled. This includes crisis detection and privacy protections required for the mode.

4.2 Safety tier

Safety tier controls how deeply SASI analyzes each message.

Typical tiers

1. **Fast**
Designed for low risk business contexts where speed is critical
2. **Standard**
Balanced for most applications
3. **Maximum**
Designed for clinical, youth, and safety critical contexts where coverage is prioritized

Public edition note

Exact performance targets and internal phase breakdowns are provided in the customer pack.

4.3 Privacy and PII protection

PII protection controls what categories of identifiers are detected and redacted.

Typical privacy levels

1. **Minimal**
Basic financial and high risk identifiers
2. **Standard**
Common identifiers such as contact information
3. **Regulated healthcare grade**
Broad identifier coverage aligned to healthcare privacy expectations
4. **Youth maximum**
Expanded coverage suitable for child and student contexts

Safety locks

In regulated contexts, privacy levels cannot be reduced below the required baseline.

4.4 Crisis sensitivity controls

Crisis sensitivity defines how early SASI flags potential crisis content and when it escalates.

Controls you will typically manage

1. **Explicit crisis detection sensitivity**
For direct statements indicating imminent harm
2. **Passive distress sensitivity**
For indirect signals that may require monitoring
3. **Anti hair trigger protections**
Guardrails that reduce single message false escalations and require patterns over time
4. **Context based uplifts**
Safety increases when contextual risk patterns appear, such as repeated distress, time based vulnerability windows, or masking language

Public edition note

Specific thresholds, ranges, and internal decision breakpoints are provided only to customers.

4.5 Adversarial and obfuscation detection

SASI can detect attempts to bypass safety controls, including instruction injection and obfuscated text.

Common controls

1. Detection enabled
2. Sensitivity level
3. Enforcement choice
Flag only, rate limit, or block, depending on mode and policy

Regulated mode enforcement

In higher risk modes, these protections are always enabled.

4.6 Crisis rewriting and safe summaries

In certain modes, SASI can transform crisis content before it reaches the model. This prevents accidental disclosure and reduces the chance of unsafe model behavior.

What this means operationally

1. Your product receives a safe representation suitable for downstream handling
2. Your incident workflow still retains the decision context for audits
3. Your application can route the user to human review or crisis resources without passing sensitive content into the model

Public edition note

The exact templates and structured summary formats are customer only.

4.7 Notifications and integrations

SASI supports event driven notifications for safety critical events.

Common events

1. Crisis detected
2. Monitoring activated
3. Escalation required

4. System health changes
5. Drift alerts

Common destinations

1. Your backend
2. Incident management tools
3. On call paging workflows
4. Clinical team routing systems

Public edition note

Exact payload structures and event schemas are provided to customers.

4.8 Model profile tuning

Different model providers can behave differently. SASI can apply profile tuning to keep safety behavior consistent across vendors.

When this matters

1. You switch model vendors
2. Vendor updates change refusal behavior
3. You see drift in crisis sensitivity or boundaries
4. You run multiple models and want consistent governance outputs

5. Reports and Analytics

Where to find reports

Reports are generally available in the Reports section of the console, grouped by purpose.

Common report types

1. Audit logs
2. Session trails
3. Crisis and escalation reports
4. Privacy and redaction reports
5. Performance and reliability reports
6. Compliance exports

5.1 Audit logs

Audit logs record the decision outcome for each analyzed message.

Typical fields included

1. Timestamp
2. Mode and policy context
3. Risk level and recommended action
4. Safety and emotional signal summaries
5. Privacy detection types, never the sensitive values
6. Decision explanation fields suitable for auditors
7. Processing metadata and health markers

Typical fields excluded

1. Raw user message text in privacy protected deployments
2. Raw sensitive identifiers

3. Generated model responses

Immutability and integrity

Audit records are designed to be tamper evident. Administrative access is logged.

5.2 Session audit trails

Session trails provide an audit view across a full interaction sequence, including escalations, monitoring activation, and policy decisions over time.

Common uses

1. Incident review
2. Quality assurance
3. Regulatory inspection readiness
4. Clinical supervision workflows

5.3 Crisis reports

Crisis reports focus on crisis class events, escalations, and resolution tracking.

Common uses

1. Safety operations reviews
2. Clinical quality reporting
3. Trend analysis and policy tuning
4. Evidence for governance stakeholders

5.4 Privacy reports

Privacy reports summarize what types of identifiers are being detected and redacted.

Important note

Reports track categories and counts, not the sensitive values.

5.5 Performance reports

Performance reports show latency, throughput, errors, and reliability indicators.

Common uses

1. SLA monitoring
2. Capacity planning
3. Integration troubleshooting
4. Operational readiness

5.6 Compliance exports

Compliance exports produce regulator oriented packages for governance and audit workflows.

Exports typically include

1. Audit records and session trails for the selected scope
2. Integrity markers for tamper evidence
3. Configuration snapshots relevant to the audit period
4. Access logs for sensitive views, depending on mode

Public edition note

Exact export formats and field level mappings are customer only.

6. Governance, Audits, and Change Management

External audits

When preparing for an external audit, teams typically do the following.

1. Gather exports for the audit period
2. Provide configuration documentation and change history
3. Verify integrity markers and retention behavior
4. Provide evidence of role based access control and access logging
5. Produce incident summaries for the period if requested

Incident reviews

A standard incident review workflow usually includes:

1. Pull the session trail and supporting audit entries
2. Review the decision explanation and policy context
3. Confirm escalation handling in your product
4. Classify the event outcome
5. Document learnings and any tuning actions
6. Track changes through formal change control

Compliance checks

Suggested cadence

1. Weekly
Confirm audit generation, notification delivery, and health indicators
2. Monthly
Review false positive indicators, escalation performance, and privacy metrics
3. Quarterly
Run a full audit readiness drill including export validation and access review
4. Annually
Full governance review, retention confirmation, and policy refresh

Access controls

SASI supports role based access. A common model includes:

1. Administrator
Can configure policies and export compliance packages, with approvals for safety critical changes
2. Compliance officer
Read access to settings, export capabilities, and audit visibility
3. Clinical or safety staff
Access to incidents and session trails appropriate to their scope
4. Operations staff
Access to health and performance views without broad audit visibility in privacy protected deployments

7. Performance Monitoring

Key indicators

1. Crisis detection rate by mode
2. Escalation rate and response time
3. False positive indicators

4. Privacy detection rate by mode
5. Processing time and error rate
6. Adversarial attempt rate

Real time monitoring

Monitoring views typically provide:

1. Live processing health
2. Active escalations and monitoring sessions
3. Alert status and delivery health
4. Short term and long term trends

Drift monitoring

SASI can monitor for behavior drift over time, especially when model vendors update systems silently.

What drift monitoring looks for

1. Changes in crisis sensitivity
2. Changes in refusal behavior patterns
3. Boundary violations against the active mode policy
4. Degradation in decision consistency

When to review

1. Weekly as a baseline
2. Immediately after model vendor changes
3. After major policy changes in your product

What teams typically do when drift is detected

1. Review the drift report
2. Confirm whether a vendor update occurred
3. Apply model profile adjustments if applicable
4. Review with safety or clinical leadership for high risk modes
5. Document actions taken in the change log

8. Operational Modes

SASI uses multiple operational modes grouped by risk category. Modes differ in:

1. Crisis sensitivity posture
2. Privacy scope
3. Audit depth and retention expectations
4. Available governance features
5. Required partner responsibilities

Common mode categories

1. Highest sensitivity
Clinical and youth contexts
2. Elevated sensitivity
Mental wellness and education contexts
3. Standard sensitivity
General consumer and coaching contexts

4. Lower sensitivity
Business contexts with reduced crisis relevance

Partner responsibilities by mode

In regulated modes, your application is expected to implement certain workflows, such as:

1. Escalation routing to humans where applicable
2. Age gating and consent handling in child contexts
3. Secure retention behavior aligned to your compliance posture
4. Use of privacy protected message representations when required
5. Access control and access logging for sensitive views

Public edition note

Exact mode names, default policy values, and required integration flags are customer only.

9. Quick Reference

Risk levels

Level	Meaning	Typical product response
Safe	No safety concern	Continue
Low	Minor concern	Continue
Moderate	Noticeable distress or risk	Add empathy or monitor
Elevated	Stronger concern	Monitor or escalate
High	High concern	Escalate
Critical	Immediate safety concern	Present crisis resources and escalate

Common recommended actions

Action	When used	Typical handling
Continue	Normal messages	Proceed
Empathy	Distress present	Respond with care and grounding
Monitor	Ongoing concern	Activate monitoring workflow
Escalate	High risk	Route to human review
Crisis resources	Critical risk	Present crisis resources immediately
Block or contain	Adversarial or harmful	Block, rate limit, or isolate

Safety tiers

Tier	Typical use	Notes
Fast	Low risk business	Prioritizes speed
Standard	Most apps	Default recommendation
Maximum	Clinical and youth	Prioritizes coverage

Privacy levels

Level	Typical use	Notes
Minimal	Business workflows	Basic identifier coverage
Standard	General apps	Common identifiers
Regulated healthcare grade	Healthcare	Broad identifiers
Youth maximum	Child and student	Expanded coverage

Retention overview

Retention is policy driven and depends on your mode and compliance posture. Regulated contexts generally require longer retention and stricter access logging.

10. Support and Resources

Documentation available to customers

Customers receive an implementation pack that includes:

1. Integration guide with exact field mapping
2. API and event schema reference
3. Audit and export specification
4. Compliance mapping notes and configuration snapshot guidance

Getting help

1. **Technical issues**
Contact SASI support or your assigned implementation partner
2. **Configuration questions**
For safety critical changes, involve your clinical or safety lead
3. **Compliance questions**
Involve your compliance officer or legal counsel

Change management best practice

Before changes

1. Record the current configuration snapshot
2. Obtain approvals for safety critical changes
3. Test in a non production environment
4. Document rationale and expected impact

After changes

1. Monitor metrics for drift or unexpected changes
2. Review incident rates and false positive indicators
3. Document outcomes and close the change record